

# Credit Risk Resampling Techniques

In [30]:

```
import warnings
warnings.filterwarnings('ignore')
```

In [31]:

```
import numpy as np
import pandas as pd
from pathlib import Path
from collections import Counter
```

## Read the CSV and Perform Basic Data Cleaning

In [32]:

```
columns = [
    "loan_amnt", "int_rate", "installment", "home_ownership",
    "annual_inc", "verification_status", "issue_d", "loan_status",
    "pymnt_plan", "dti", "delinq_2yrs", "inq_last_6mths",
    "open_acc", "pub_rec", "revol_bal", "total_acc",
    "initial_list_status", "out_prncp", "out_prncp_inv", "total_pymnt",
    "total_pymnt_inv", "total_rec_prncp", "total_rec_int", "total_rec_late_fee",
    "recoveries", "collection_recovery_fee", "last_pymnt_amnt", "next_pymnt_d",
    "collections_12_mths_ex_med", "policy_code", "application_type", "acc_now_de
linq",
    "tot_coll_amt", "tot_cur_bal", "open_acc_6m", "open_act_il",
    "open_il_12m", "open_il_24m", "mths_since_rcnt_il", "total_bal_il",
    "il_util", "open_rv_12m", "open_rv_24m", "max_bal_bc",
    "all_util", "total_rev_hi_lim", "inq_fi", "total_cu_tl",
    "inq_last_12m", "acc_open_past_24mths", "avg_cur_bal", "bc_open_to_buy",
    "bc_util", "chargeoff_within_12_mths", "delinq_amnt", "mo_sin_old_il_acct",
    "mo_sin_old_rev_tl_op", "mo_sin_rcnt_rev_tl_op", "mo_sin_rcnt_tl", "mort_acc
",
    "mths_since_recent_bc", "mths_since_recent_inq", "num_accts_ever_120_pd", "n
um_actv_bc_tl",
    "num_actv_rev_tl", "num_bc_sats", "num_bc_tl", "num_il_tl",
    "num_op_rev_tl", "num_rev_accts", "num_rev_tl_bal_gt_0",
    "num_sats", "num_tl_120dpd_2m", "num_tl_30dpd", "num_tl_90g_dpd_24m",
    "num_tl_op_past_12m", "pct_tl_nvr_dlq", "percent_bc_gt_75", "pub_rec_bankrup
tcies",
    "tax_liens", "tot_hi_cred_lim", "total_bal_ex_mort", "total_bc_limit",
    "total_il_high_credit_limit", "hardship_flag", "debt_settlement_flag"
]

target = ["loan_status"]
```

In [33]:

```
# Load the data
file_path = Path('../Resources/LoanStats_2019Q1.csv.zip')
df = pd.read_csv(file_path, skiprows=1)[: -2]
df = df.loc[:, columns].copy()

# Drop the null columns where all values are null
df = df.dropna(axis='columns', how='all')

# Drop the null rows
df = df.dropna()

# Remove the `Issued` loan status
issued_mask = df['loan_status'] != 'Issued'
df = df.loc[issued_mask]

# convert interest rate to numerical
df['int_rate'] = df['int_rate'].str.replace('%', '')
df['int_rate'] = df['int_rate'].astype('float') / 100

# Convert the target column values to low_risk and high_risk based on their values
x = {'Current': 'low_risk'}
df = df.replace(x)

x = dict.fromkeys(['Late (31-120 days)', 'Late (16-30 days)', 'Default', 'In Grace Period'], 'high_risk')
df = df.replace(x)

df.reset_index(inplace=True, drop=True)

df.head()
```

Out[33]:

	loan_amnt	int_rate	installment	home_ownership	annual_inc	verification_status	issue_d	I
0	10500.0	0.1719	375.35	RENT	66000.0	Source Verified	Mar-2019	
1	25000.0	0.2000	929.09	MORTGAGE	105000.0	Verified	Mar-2019	
2	20000.0	0.2000	529.88	MORTGAGE	56000.0	Verified	Mar-2019	
3	10000.0	0.1640	353.55	RENT	92000.0	Verified	Mar-2019	
4	22000.0	0.1474	520.39	MORTGAGE	52000.0	Not Verified	Mar-2019	

5 rows × 86 columns

## Split the Data into Training and Testing

In [34]:

```
# Create our features
X = pd.get_dummies(df.drop(columns='loan_status', axis = 'columns'))# YOUR CODE HERE
```

In [35]:

```
X.describe()
```

Out[35]:

	loan_amnt	int_rate	installment	annual_inc	dti	delinq_2yrs	i
<b>count</b>	68817.000000	68817.000000	68817.000000	6.881700e+04	68817.000000	68817.000000	
<b>mean</b>	16677.594562	0.127718	480.652863	8.821371e+04	21.778153	0.217766	
<b>std</b>	10277.348590	0.048130	288.062432	1.155800e+05	20.199244	0.718367	
<b>min</b>	1000.000000	0.060000	30.890000	4.000000e+01	0.000000	0.000000	
<b>25%</b>	9000.000000	0.088100	265.730000	5.000000e+04	13.890000	0.000000	
<b>50%</b>	15000.000000	0.118000	404.560000	7.300000e+04	19.760000	0.000000	
<b>75%</b>	24000.000000	0.155700	648.100000	1.040000e+05	26.660000	0.000000	
<b>max</b>	40000.000000	0.308400	1676.230000	8.797500e+06	999.000000	18.000000	

8 rows × 95 columns

In [36]:

```
# Create our target
y = df['loan_status']# YOUR CODE HERE
```

In [37]:

```
y.value_counts()
```

Out[37]:

```
low_risk      68470
high_risk      347
Name: loan_status, dtype: int64
```

In [38]:

```
# Create X_train, X_test, y_train, y_test
# YOUR CODE HERE
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test= train_test_split(X,
                                                    y,
                                                    random_state=1
                                                    )

X_train.shape
```

Out[38]:

(51612, 95)

In [39]:

```
Counter(y_train)
```

Out[39]:

Counter({'low\_risk': 51366, 'high\_risk': 246})

## Oversampling

In this section, you will compare two oversampling algorithms to determine which algorithm results in the best performance. You will oversample the data using the naive random oversampling algorithm and the SMOTE algorithm. For each algorithm, be sure to complete the following steps:

1. View the count of the target classes using `Counter` from the collections library.
2. Use the resampled data to train a logistic regression model.
3. Calculate the balanced accuracy score from `sklearn.metrics`.
4. Print the confusion matrix from `sklearn.metrics`.
5. Generate a classification report using the `imbalanced_classification_report` from `imbalanced-learn`.

Note: Use a random state of 1 for each sampling algorithm to ensure consistency between tests

## Naive Random Oversampling

In [40]:

```
# Resample the training data with the RandomOverSampler
# YOUR CODE HERE
from imblearn.over_sampling import RandomOverSampler

ros = RandomOverSampler(random_state=1)
X_resampled, y_resampled = ros.fit_resample(X_train, y_train)
Counter(y_resampled)
```

Out[40]:

```
Counter({'low_risk': 51366, 'high_risk': 51366})
```

In [44]:

```
# Logistic regression using random oversampled data
from sklearn.linear_model import LogisticRegression

model = LogisticRegression(solver='liblinear', random_state=1)
model.fit(X_resampled, y_resampled)
```

Out[44]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=1, solver='liblinear', tol=0.0001, verbose=0,
                    warm_start=False)
```

In [45]:

```
y_pred = model.predict(X_test)
```

In [46]:

```
from sklearn.metrics import balanced_accuracy_score

balanced_accuracy_score(y_test, y_pred)
```

Out[46]:

```
0.7163908158823367
```

In [47]:

```
# Display the confusion matrix
from sklearn.metrics import confusion_matrix

confusion_matrix(y_test, y_pred)
```

Out[47]:

```
array([[ 73,  28],
       [4960, 12144]])
```

In [48]:

```
from imblearn.metrics import classification_report_imbalanced

print(classification_report_imbalanced(y_test, y_pred))
```

		pre	rec	spe	f1	geo
iba	sup					
high_risk		0.01	0.72	0.71	0.03	0.72
0.51	101					
low_risk		1.00	0.71	0.72	0.83	0.72
0.51	17104					
avg / total		0.99	0.71	0.72	0.82	0.72
0.51	17205					

## SMOTE Oversampling

In [50]:

```
# Resample the training data with SMOTE
# YOUR CODE HERE
from imblearn.over_sampling import SMOTE
X_resampled, y_resampled = SMOTE(random_state=1, sampling_strategy=1.0).fit_resample(
    X_train, y_train
)
from collections import Counter
Counter(y_resampled)
```

Out[50]:

```
Counter({'low_risk': 51366, 'high_risk': 51366})
```



In [74]:

```
# Train the Logistic Regression model using the resampled data
# YOUR CODE HERE
model = LogisticRegression(solver = 'liblinear', random_state=1)
model.fit(X_resampled, y_resampled)
```

Out[74]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=1, solver='liblinear', tol=0.0001, verbose=0,
                    warm_start=False)
```

In [75]:

```
y_pred = model.predict(X_test)
```

In [76]:

```
# Calculated the balanced accuracy score
# YOUR CODE HERE
balanced_accuracy_score(y_test, y_pred)
```

Out[76]:

```
0.7107375728218286
```

In [77]:

```
# Display the confusion matrix
# YOUR CODE HERE
confusion_matrix(y_test, y_pred)
```

Out[77]:

```
array([[ 69,  32],
       [4476, 12628]])
```

In [78]:

```
# Print the imbalanced classification report
# YOUR CODE HERE
print(classification_report_imbalanced(y_test, y_pred))
```

iba	sup	pre	rec	spe	f1	geo
high_risk	0.50	0.02	0.68	0.74	0.03	0.71
low_risk	0.51	1.00	0.74	0.68	0.85	0.71
avg / total	0.51	0.99	0.74	0.68	0.84	0.71

## Undersampling

In this section, you will test an undersampling algorithms to determine which algorithm results in the best performance compared to the oversampling algorithms above. You will undersample the data using the Cluster Centroids algorithm and complete the following steps:

1. View the count of the target classes using `Counter` from the collections library.
2. Use the resampled data to train a logistic regression model.
3. Calculate the balanced accuracy score from `sklearn.metrics`.
4. Print the confusion matrix from `sklearn.metrics`.
5. Generate a classification report using the `imbalanced_classification_report` from `imbalanced-learn`.

Note: Use a random state of 1 for each sampling algorithm to ensure consistency between tests

In [79]:

```
# Resample the data using the ClusterCentroids resampler
# YOUR CODE HERE
from imblearn.under_sampling import ClusterCentroids

cc = ClusterCentroids(random_state=1)
X_resampled, y_resampled = cc.fit_resample(X_train, y_train)

from collections import Counter

Counter(y_resampled)
```

Out[79]:

```
Counter({'high_risk': 246, 'low_risk': 246})
```

In [91]:

```
# Train the Logistic Regression model using the resampled data
# YOUR CODE HERE
model = LogisticRegression(solver='liblinear', random_state=1)
model.fit(X_resampled, y_resampled)
```

Out[91]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=1, solver='liblinear', tol=0.0001, verbose=0,
                    warm_start=False)
```

In [92]:

```
y_pred = model.predict(X_test)
balanced_accuracy_score(y_test, y_pred)
```

Out[92]:

```
0.6564285813520547
```

In [93]:

```
# Display the confusion matrix
# YOUR CODE HERE
confusion_matrix(y_test, y_pred)
```

Out[93]:

```
array([[ 84,  17],
       [8874, 8230]])
```

In [94]:

```
# Print the imbalanced classification report
# YOUR CODE HERE
print(classification_report_imbalanced(y_test, y_pred))
```

		pre	rec	spe	f1	geo
iba	sup					
high_risk		0.01	0.83	0.48	0.02	0.63
0.41	101					
low_risk		1.00	0.48	0.83	0.65	0.63
0.39	17104					
avg / total		0.99	0.48	0.83	0.65	0.63
0.39	17205					

## Combination (Over and Under) Sampling

In this section, you will test a combination over- and under-sampling algorithm to determine if the algorithm results in the best performance compared to the other sampling algorithms above. You will resample the data using the SMOTEENN algorithm and complete the following steps:

1. View the count of the target classes using `Counter` from the `collections` library.
2. Use the resampled data to train a logistic regression model.
3. Calculate the balanced accuracy score from `sklearn.metrics`.
4. Print the confusion matrix from `sklearn.metrics`.
5. Generate a classification report using the `imbalanced_classification_report` from `imbalanced-learn`.

Note: Use a random state of 1 for each sampling algorithm to ensure consistency between tests

In [100]:

```
# Resample the training data with SMOTEENN
# YOUR CODE HERE
from imblearn.combine import SMOTEENN

smote_enn = SMOTEENN(random_state=0)
X_resampled, y_resampled = smote_enn.fit_resample(X, y)
Counter(y_resampled)
```

Out[100]:

```
Counter({'high_risk': 68460, 'low_risk': 62011})
```

In [101]:

```
# Train the Logistic Regression model using the resampled data
# YOUR CODE HERE
model = LogisticRegression(solver='liblinear', random_state=1)
model.fit(X_resampled, y_resampled)
```

Out[101]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=1, solver='liblinear', tol=0.0001, verbose=0,
                    warm_start=False)
```

In [102]:

```
# Calculated the balanced accuracy score
# YOUR CODE HERE
y_pred = model.predict(X_test)
balanced_accuracy_score(y_test, y_pred)
```

Out[102]:

```
0.7317583635117487
```

In [103]:

```
# Display the confusion matrix
# YOUR CODE HERE
confusion_matrix(y_test, y_pred)
```

Out[103]:

```
array([[ 75,  26],
       [4773, 12331]])
```

In [104]:

```
# Print the imbalanced classification report
# YOUR CODE HERE
print(classification_report_imbalanced(y_test, y_pred))
```

		pre	rec	spe	f1	geo
iba	sup					
high_risk		0.02	0.74	0.72	0.03	0.73
0.54	101					
low_risk		1.00	0.72	0.74	0.84	0.73
0.53	17104					
avg / total		0.99	0.72	0.74	0.83	0.73
0.53	17205					